

Assignment 2 Image captioning Digits

Group-71

Name : Saurabh Arunrao Dhande - 2021FC04700

```
In [34]: # Import necessary Libraries
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import matplotlib.pyplot as plt
import numpy as np
```

Data Procession

```
In [2]: # Load MNIST dataset
tf.keras.datasets.mnist.load_data(path="mnist.npz")
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
 11490434/11490434 [=====] - 1s 0us/step
 (60000, 28, 28)
 (60000,)
 (10000, 28, 28)
 (10000,)


```
In [4]: # Generate random indices for two samples
sample_1, sample_2 = np.random.randint(0, len(x_train), size=2)

# Plot the first sample
plt.figure()
plt.imshow(x_train[sample_1], cmap='gray')
plt.title("caption_digit: " + str(y_train[sample_1]))
plt.axis('off')

# Plot the second sample
plt.figure()
plt.imshow(x_train[sample_2], cmap='gray')
plt.title("caption_digit: " + str(y_train[sample_2]))
plt.axis('off')

plt.show()
```

caption_digit: 2



caption_digit: 3



```
In [5]: #updating replacing inter with string of number in y_train and y_test
int_to_str = {
    0: 'zero',1: 'one',2: 'two',3: 'three',4: 'four',5: 'five',6: 'six',7: 'seven',8: 'eig
}

str_to_int = {v: k for k, v in int_to_str.items()} # create new dictionary to map string

y_test = [int_to_str[num] if isinstance(num, int) else num for num in y_test]
y_train = [int_to_str[num] if isinstance(num, int) else num for num in y_train]

y_test = [str_to_int[label] if label in str_to_int else label for label in y_test] # conv
y_train = [str_to_int[label] if label in str_to_int else label for label in y_train] # cc

y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)

print(y_train)
print(y_test)
```

```
[[0. 0. 0. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 ...
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 1. 0.]]
[[0. 0. 0. ... 1. 0. 0.]
 [0. 0. 1. ... 0. 0. 0.]
 [0. 1. 0. ... 0. 0. 0.]
 ...
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]]
```

```
In [6]: # Reshape and normalize input data
x_train = x_train.reshape((-1, 28, 28, 1))
x_test = x_test.reshape((-1, 28, 28, 1))
x_train = x_train / 255.0
x_test = x_test / 255.0
```

Model Building

layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(28, 28, 1)): This layer applies 32 filters of size 3x3 to the input image of size 28x28x1, and uses the ReLU activation function. The choice of 32 filters is a common default in many CNN architectures, and the use of ReLU activation helps introduce non-linearity into the model.

layers.MaxPooling2D(pool_size=(2, 2)): This layer performs max pooling on the output of the previous convolutional layer, reducing the spatial dimensions of the output by a factor of 2 in each dimension. This helps reduce the number of parameters in the model and introduces some degree of translation invariance.

layers.Reshape(target_shape=(5408,)): This layer reshapes the output of the previous layer into a 1D array of size 5408. This is required to feed the output to the next LSTM layer.

layers.RepeatVector(28): This layer repeats the output of the previous layer 28 times, which is necessary for feeding the data to the LSTM layer in a sequence format.

layers.LSTM(units=64, return_sequences=True): This layer applies an LSTM layer with 64 units and returns the output sequence. LSTMs are commonly used in sequence modeling tasks due to their ability to capture long-term dependencies.

layers.Flatten(): This layer flattens the output of the LSTM layer into a 1D array.

layers.Dense(units=10, activation='softmax'): This layer applies a dense layer with 10 units and a softmax activation function. The softmax function is commonly used for multi-class classification tasks to output probability scores for each class.

```
In [27]: # Define CNN+LSTM model architecture
model = keras.Sequential([
    layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(28, 28,
    layers.MaxPooling2D(pool_size=(2, 2)),
    layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),
    layers.MaxPooling2D(pool_size=(2, 2)),
```

```

layers.Reshape(target_shape=(1600,)),
layers.RepeatVector(28),
layers.LSTM(units=64, return_sequences=True),
layers.Flatten(),
layers.Dense(units=10, activation='softmax')
])

model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_2 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_2 (MaxPooling 2D)	(None, 13, 13, 32)	0
conv2d_3 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_3 (MaxPooling 2D)	(None, 5, 5, 64)	0
reshape_1 (Reshape)	(None, 1600)	0
repeat_vector_1 (RepeatVector)	(None, 28, 1600)	0
lstm_1 (LSTM)	(None, 28, 64)	426240
flatten_1 (Flatten)	(None, 1792)	0
dense_1 (Dense)	(None, 10)	17930
=====		
Total params: 462,986		
Trainable params: 462,986		
Non-trainable params: 0		

Model Compilation

The choice of learning rate, optimizer, and loss function in a deep learning model plays a critical role in determining its performance. In the given code snippet, the Adam optimizer with a default learning rate of 0.001 is used, which is a commonly used optimizer for deep learning models. The choice of the categorical cross-entropy loss function is appropriate for multi-class classification problems. However, selecting the optimal values of these hyperparameters can be a challenging task and may require experimenting with different values. It is important to evaluate the performance of the model on a validation set to determine the best configuration for these hyperparameters to achieve optimal performance.

```

In [28]: # Compile model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

```

Model Training

```

In [29]: # Train model
history = model.fit(x_train, y_train, epochs=1, batch_size=32, validation_data=(x_test, y_



```

1875/1875 [=====] - 278s 146ms/step - loss: 0.1147 - accuracy: 0.9638 - val_loss: 0.0438 - val_accuracy: 0.9868

Note : I used epoch = 1 while fitting model because when i am trying to use epoch > 1 then google colab is getting crashed, because of this I was not able to plot loss and accuracy history graphs for both train and validation set.

```
In [30]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend(['Train', 'Validation'], loc='upper right')
plt.show()

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(['Train', 'Validation'], loc='lower right')
plt.show()
```



Model Evaluation

```
In [31]: # Evaluate model on test set
test_loss, test_acc = model.evaluate(x_test, y_test)
print('Test accuracy:', test_acc)
```

```
313/313 [=====] - 13s 41ms/step - loss: 0.0438 - accuracy: 0.9868
Test accuracy: 0.9868000149726868
```

```
In [32]: import numpy as np
from sklearn.metrics import multilabel_confusion_matrix

y_pred = model.predict(x_test)
y_pred_binary = np.round(y_pred)

cm = multilabel_confusion_matrix(y_test, y_pred_binary)

print("Confusion matrix:")
print(cm)
```

```
313/313 [=====] - 14s 42ms/step
Confusion matrix:
[[[9016   4]
  [ 12 968]]

 [[8853  12]
  [   1 1134]]

 [[8920  48]
  [   0 1032]]

 [[8989   1]
  [  21 989]]

 [[9014   4]
  [  22 960]]

 [[9104   4]
  [  10 882]]

 [[9034   8]
  [   8 950]]

 [[8962  10]
  [  27 1001]]

 [[9017   9]
  [  26 948]]

 [[8975  16]
  [  20 989]]]
```

Take 5 random images from Google and generate caption for that image.

```
In [49]: import os
import cv2
import numpy as np
import tensorflow as tf

path = '/content/sample_data/Image'
data = []
labels = []
```

```

for filename in os.listdir(path):
    if filename.endswith('.png'):
        img = cv2.imread(os.path.join(path, filename))
        img = cv2.resize(img, (28, 28))
        img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
        label = filename
        data.append(img)
        labels.append(label)

data = np.array(data)

labels = np.array(labels)
print(len(data))

data = np.expand_dims(data, axis=-1)

pred_t = model.predict(data)

```

```

5
1/1 [=====] - 0s 35ms/step

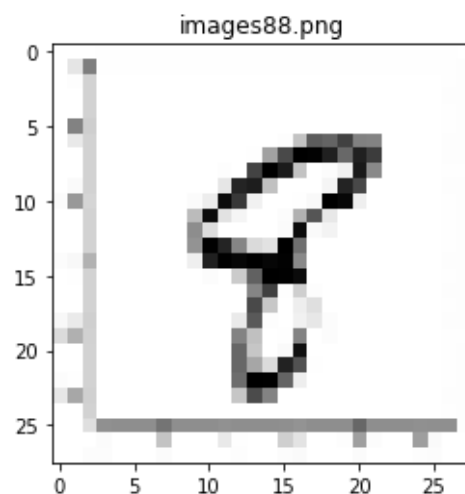

```

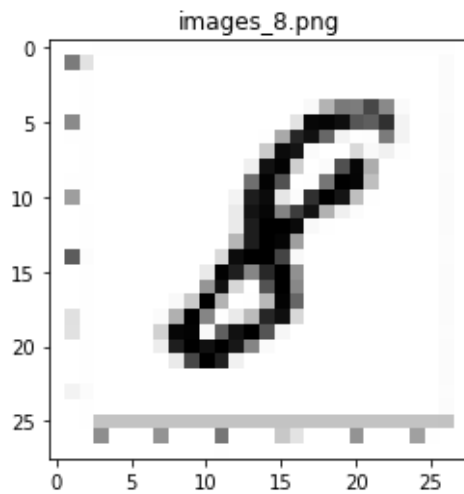
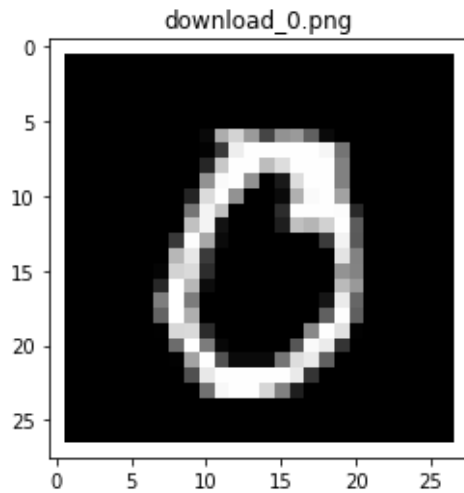
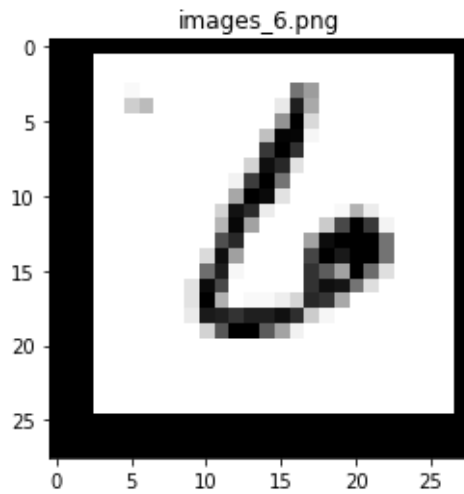
```

In [53]: import matplotlib.pyplot as plt

for i in range(len(data)):
    plt.imshow(data[i], cmap='gray')
    plt.title(labels[i])
    plt.show()

```





In [50]: `print(pred_t)`

```
[[2.3464945e-05 3.8579132e-03 2.9854311e-02 9.0560770e-01 1.1511862e-05
 8.1303615e-06 1.5556642e-06 6.0342986e-02 2.8964618e-05 2.6339572e-04]
 [2.1251183e-04 1.8848201e-03 1.4081442e-03 1.1803533e-03 1.1841732e-04
 1.7691224e-03 6.3419784e-03 5.0488070e-05 9.8580188e-01 1.2323017e-03]
 [6.7053122e-08 1.9968436e-06 3.6050689e-05 1.8143284e-07 3.8163711e-05
 1.3063962e-04 9.9945420e-01 1.1589938e-07 3.3847100e-04 6.9727669e-08]
 [9.8742861e-01 5.1744114e-06 7.3458930e-03 4.8356902e-05 2.0062951e-03
 3.9967944e-04 1.1707354e-04 1.9050970e-03 8.4908061e-06 7.3532667e-04]
 [8.1335020e-04 2.5657902e-04 3.0567318e-01 3.4225578e-04 6.6158280e-02
 2.0559282e-05 7.3542749e-04 1.2170168e-02 6.1160153e-01 2.2287061e-03]]
```



```
In [51]: for a in range(0, len(pred_t)):
          for i in range(0, len(pred_t[a])):
              #print(a[i])
              if pred_t[a][i] > 0.55:
                  output_str = ""
                  for char in int_to_str[i]:
                      output_str += ' ' + char + ', '
                  # Remove the extra space and comma at the end of the string
                  output_str = output_str[:-2]
                  print(output_str)
```

```
"t", "h", "r", "e", "e"
"e", "i", "g", "h", "t"
"s", "i", "x"
"z", "e", "r", "o"
"e", "i", "g", "h", "t"
```